


Web search for a live GTM chat agent

You.com, Tavily, Exa and Parallel, compared live on 10 queries for GTM agent, the most common use cases for web search APIs.

 Pallavi Singla

To: VP of Engineering, Reo.dev · **Re:** which web search API should power a live GTM chat agent



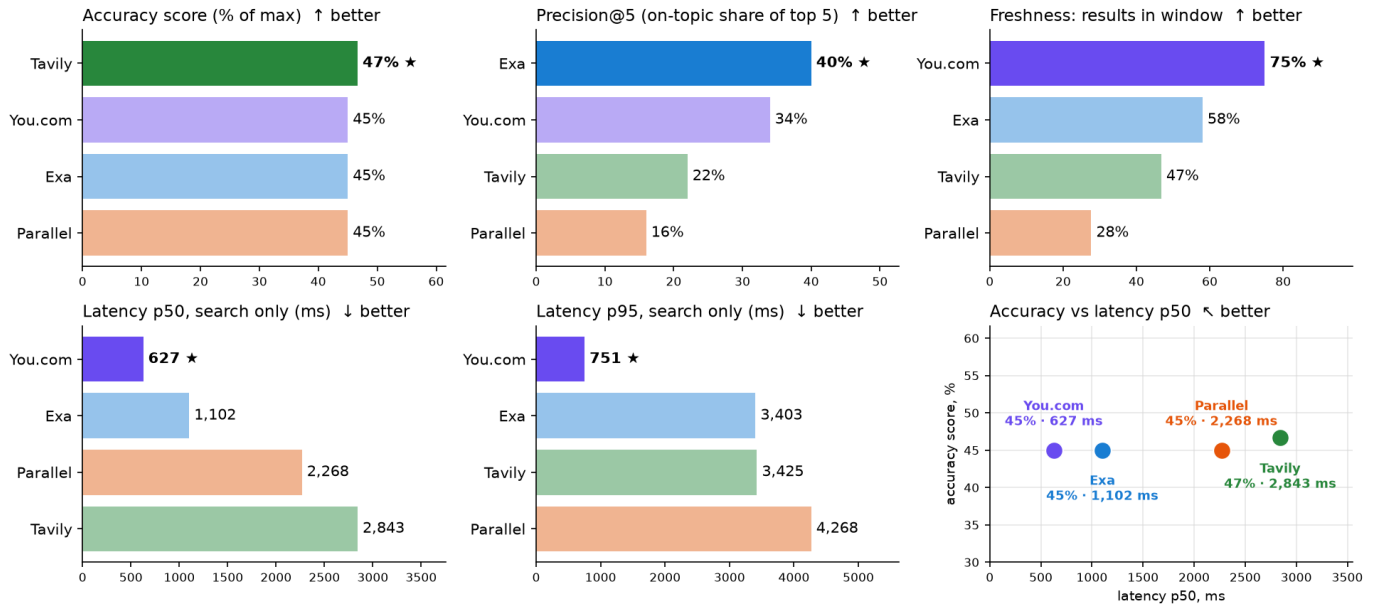
The use case: a real-time GTM chat agent

For GTM agents that answer "what's new?" in a live conversation, You.com is the web search that is both fast enough for chat and fresh enough to trust. Its accuracy matched the other three.

- 1. **Fast enough for chat:** 10 of 10 searches under 1 s; no other API managed more than 3 of 10.
- 2. **The freshest signals:** 75% of results inside the requested window, and the only API to surface today's Apache Iggy activity.
- 3. **Never comes back empty:** results on all 30 calls, against 9 empty for Tavily and 3 for Exa.

Parallel and Tavily did better for questions that need an official page, such as release notes or product docs.

GTM — all providers, ranked best to worst per metric



Bars sorted best → worst; ★ = best. Latency = client wall-clock, search only, run 1 (first time each query sent). Freshness = share of results dated inside the query's window. 10 queries per use case.

Quick glance on queries Tested:

A sales rep or founder at company called "LaserData - selling to Apache Iggy developers" asks the agent in chat, for example in Slack or the product, "What's new at Confluent (competitor) today?" or "Who's talking about Apache Iggy?"

GTM — what each API returned, per question

Each cell: the result the answer relied on (else the top result) exactly as returned — domain · published date, then title and its rank. Grey line = grader's verdict.

Colour = accuracy across 3 repeats (search + content run):

correct in all 3 repeats

partly correct, or correct in only some repeats

wrong or missed

Where the true answer is NONE (nothing happened), an empty result counts as correct.

			You.com	Tavily	Exa	Parallel
	What the rep asks	Correct answer (key)	Freshness (results in window): 75% Accuracy score: 45%	Freshness (results in window): 47% Accuracy score: 47%	Freshness (results in window): 58% Accuracy score: 45%	Freshness (results in window): 28% Accuracy score: 45%
G1	Who is hitting Kafka latency problems right now?	Kafka lag issue opened in last 6h	github.com · 2026-10-07 "[Blueprint] kafka-consumer-lag-sentinel - PR..." #1/10 why: Issues from Oct 7 (outside 6h)	0 results returned why: Empty — no results	github.com · 2026-10-08 "Issues · NousResearch/hermes-agent · GitHub" #1/9 why: Unrelated repo issue-list pages	github.com · no date "GitHub - vbforge/kafka-lag-lab: A Kafka consumer..." #1/9 why: Undated Kafka lag tool repos
		Oct 8 ATO interview, Confluent VP	confluent.safebase.us · 2026-10-07 "Confluent Trust Center Powered by SafeBase" #1/10 why: Keyword noise, no Confluent news	allthingsopen.org · 2026-10-08 "You don't have an AI strategy without a data" #5/8 why: Hit in 2 of 3 runs (ATO Oct 8)	allthingsopen.org · no date "You don't have an AI strategy without a data" #3/10 why: ATO Oct 8 interview with VP	confluent.io · no date "News, Press Releases, and Media Coverage [...]" #1/3 why: Undated newsroom, IBM deal page
G3	Any new mentions of Apache Iggy?	Iggy repo issues/PRs from last 24h	github.com · 2026-10-08 "Partition replica stuck after rolling crash..." #1/10 why: Dated apache/iggy issues (Oct 8)	reddit.com · no date "Rust talks at P99 CONF (free, virtual)" #1/8 why: Undated Reddit post, noise	thedroptimes.com · no date "News Roundup TheDropTimes" #1/10 why: Unrelated Apache pages	iggy.apache.org · no date "Hyper-Efficient Message Streaming written in Rust..." #1/4 why: Undated landing and docs pages
G4	What's in Iggy's latest release?	server-0.9.0 (Sep 18)	github.com · 2026-10-08 "Issue - apache/iggy" #1/10 why: Iggy issues, no release page	0 results returned why: Empty — no results	github.com · 2026-10-08 "Issues - apache/incubator-kie-issues - GitHub" #1/10 why: Other Apache projects' issue lists	github.com · no date "Releases - apache/iggy - GitHub" #2/6 why: server-0.9.0 release page (Sep 18)
G5	What's the official fix for consumer lag?	Official doc URL, both settings	docs.confluent.io · 2026-09-09 "Monitor Consumer Lag in Confluent Platform [...]" #1/10 why: Official docs, settings text cut off	docs.confluent.io · 2026-09-15 "Salesforce Bulk API 2.0 Sink Connector for..." #3/10 why: Confluent docs, both settings	kafka.apache.org · 2026-05-22 "Consumer and Share Consumer Configs Apache Kafka" #1/10 why: Kafka docs, only fetch.min.bytes	docs.confluent.io · no date "Consumer Metrics" #1/10 why: Docs; max.poll.records not described
G6	What's the latest NATS server release?	v2.15.1-RC.2 (Oct 6) / v2.15.0	github.com · 2026-10-05 "Run NATS Jetstream as the B.U.N.N.Y. message bus..." #1/10 why: nats-server issues, no release	github.com · 2026-10-01 "Releases - coder/coder - GitHub" #1/9 why: Other projects' release pages	github.com · 2026-10-06 "Issues - nats-io/nats-server - GitHub" #1/10 why: nats-io issue lists, no release	github.com · no date "Releases: nats-io/nats-server - GitHub" #1/9 why: v2.15.1-RC.2 and v2.15.0
G7	Is Confluent Cloud having an incident?	No active incident	status.confluent.clo... · no date "Confluent Cloud Status" #1/2 why: No incident; Sep 21 page is stale	0 results returned why: Empty — no results	0 results returned why: Empty — no results	status.confluent.clo... · no date "Incident History - Confluent Cloud Status" #1/1 why: Incident history header only
G8	Which AI-agent infrastructure startups just raised?	Agent/data-infra raise in last 48h	techstartups.com · 2026-10-08 "Startup Funding News Today, October 8, 2026..." #2/10 why: Manus \$500M+ (Oct 8)	htgf.de · 2026-10-08 "TODAY Secures €2.8M to scale compliant AI agents..." #3/9 why: TODAY €2.8M, Monid \$7.7M	dealroom.co · no date "Amperand raises \$15M to connect AI agents with..." #1/10 why: 5 in-window raises incl. Rein \$25M	dealroom.co · 2026-10-08 "Rein Security raises \$25M to secure enterprise AI..." #2/5 why: Rein \$25M Series A (Oct 8)
G9	What's new this week on memory for AI agents?	Agent-memory articles this week	arize.com · 2026-10-02 "AI memory: types, architecture, and how memory..." #3/10 why: Arize, Ability.ai, dev.to (Oct 2-6)	ability.ai · 2026-10-06 "AI agent durability: why humans are the ultimate..." #1/8 why: Ability.ai, Wenlan, Eigent (Oct 4-7)	likeone.ai · 2026-10-02 "Persistent Memory for AI Agents (2026) Like One" #1/10 why: Like One, Databricks, arXiv (Oct 2-6)	arxiv.org · no date "Memory for Autonomous LLM Agents: Mechanisms..." #1/6 why: On-topic but nothing dated in window
G10	Any Redpanda vs Kafka latency benchmarks?	In-window Redpanda vs Kafka p99 data	automq.com · 2026-05-07 "Redpanda vs Kafka vs AutoMQ: Latency & Cost..." #1/10 why: Medium p99.99 numbers, no runner	robustmq.com · 2026-03-28 "RobustMQ vs Redpanda..." #1/10 why: Vendor pages, no benchmark	computingforgeeks.com · 2026-04-08 "Kafka vs Redpanda: Real Benchmarks on Identical..." #1/10 why: 5 dated benchmarks, no p99 numbers	computingforgeeks.com · 2026-04-08 "Kafka vs Redpanda: Real Benchmarks on Identical..." #1/8 why: ComputingForGeeks, avg latency only

How to use these Search API Providers across your use-cases:

GTM use case	Best fit	Why (evidence from this test)
Fresh-signal sweeps the agent re-runs every few minutes to find out targeted developers	You.com	Freshest results, and the newest relevant result. Only provider to surface today's activity.
Live chat answers for reps ("What's new at Confluent today?, Funding and startup-news alerts")	You.com	Only provider with every search under 1 s
Release notes and changelogs for competitive briefs	Parallel	Only provider to find both release pages
Product docs and configuration answers for technical replies	Tavily	Only provider whose results described both config settings.
Thought-leadership content	Exa	Most on-topic top 5 results

What we'd test next, and what would change the call

We are open to test on your real-time use-cases to help you showcase how it can help your team or needs custom updates:

Testing at scale on your real-time use-cases	Why	How it could change the recommendation
Replay 200 real questions from your reps' chat history, scoring the full answer and turn time	Real questions can be messier to provide different metric across providers	Expected: You.com stays within a few points of the best on accuracy while staying the fastest. Would change only if its accuracy falls more than 10 points below the leader. Then use You.com as the fast first pass, and send unresolved questions to Exa (News) or Parallel (official pages).
Run the same set at peak chat volume (many concurrent users)	To extrapolate and verify latency at scale.	Expected: You.com keeps its sub-second p95. Its server time was about 0.4 s, which leaves headroom under load. Would change only if another API also stays under 1 s at peak. There might be a trade-off on accuracy and cost.
Cost at your monthly search volume	You.com cost more per call in our setup	Expected: at chat volume, fewer empty results and follow-up searches offset You.com's higher price per call. Would change if cost per useful answer still tips the balance. Then use You.com for live chat turns and a cheaper API (Parallel or Exa) for batch enrichment.

Full data, methodology and all three use cases: [4 Web Search APIs, Tested Live on 3 Agent Use Cases.](#)

Independent hands-on evaluation, run live on October 8, 2026. Not affiliated with or endorsed by any provider named.